

# Physical Database Design

## Indexes Part 3 : Clustered and non-clustered Indexes



# Clustered Index

in-disk b-tree data structure consist of 3 levels :

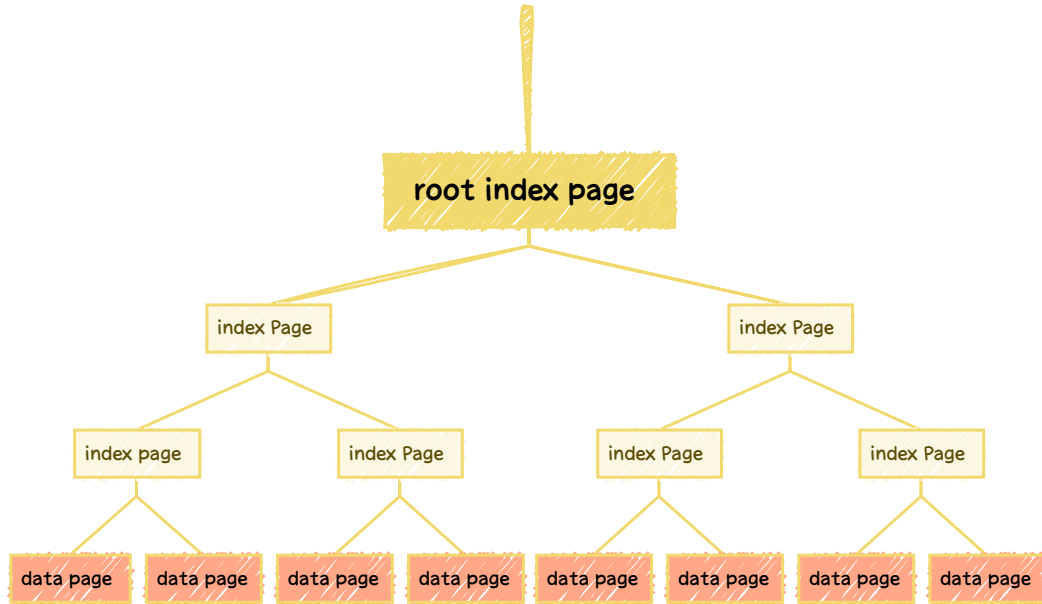
1- root node : contain  
- clustered index keys  
- with address of pages of other lower index pages

2- intermediate nodes : contain  
- clustered index keys  
- with address of pages of other lower index pages  
- or address of actual data pages

3- leaf nodes : contain  
- actual data pages

data is stored on disk with order  
- Index Key Is Sorted in Indexes Pages in Root or intermediate nodes  
- rows is Sorted ( Logically or Physically ) in Actual Data Pages in last level

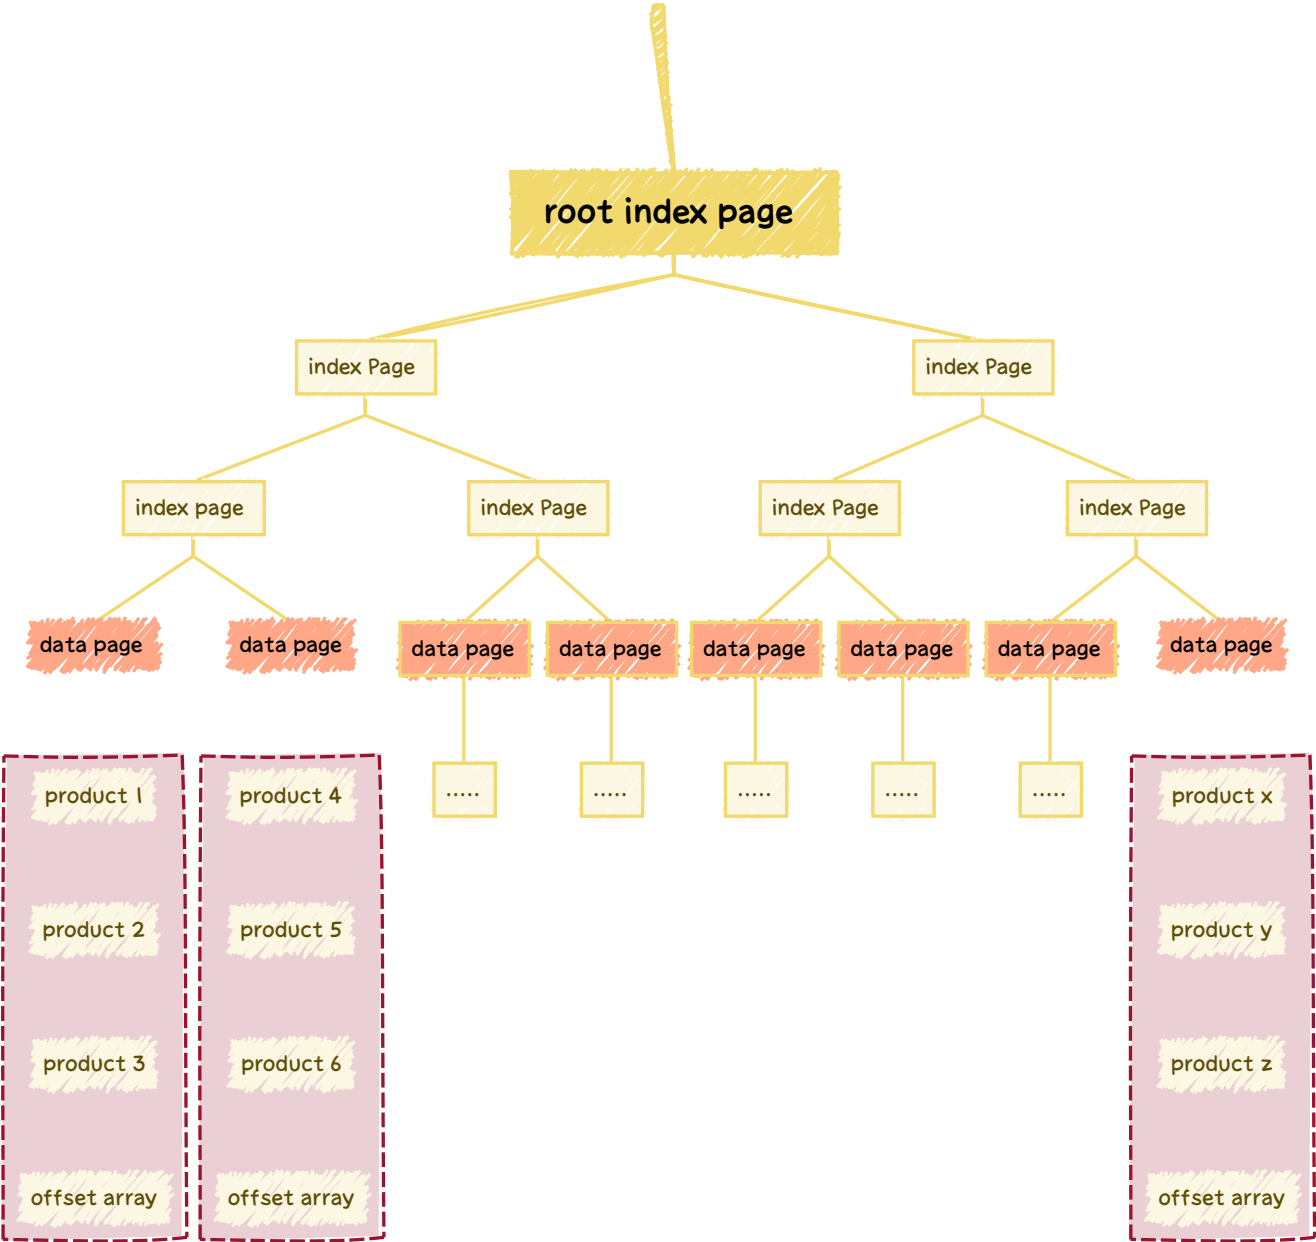
# in-disk b-tree data structure



Create Clustered Index on  
ID Column on Product Table

```
CREATE CLUSTERED INDEX [IX_Product_ID]  
ON [dbo].[Product]  
(  
    [ID] ASC  
)
```

# Product Cluster Index Table



# Clustered Index On ID Column

First Level : Root Node

page 1000	
index key	page id
30	1001
50	1002
.....	.....
z	100X

Second Level : intermediate nodes

page 1001		page 1002			
index key	page id	index key	page id	index key	page id
12	1003	35	1005	....	.....
20	1004	45	1006	....	.....
.....	.....	.....	.....	Z	.....
page 1003		page 1004			
index key	page id	index key	page id	index key	page id
3	7001	.....	.....	.....	.....
6	7002	.....	.....	.....	.....
.....	.....	.....	.....	.....	.....

Third Level : Leaf Nodes

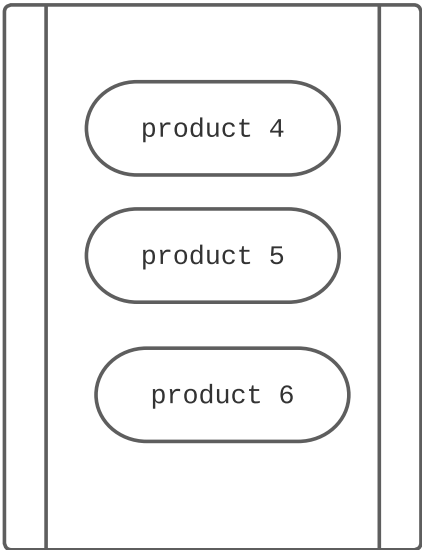
page 7001				page 7002				page 7003							
ID	name	price	Qty	ID	name	price	Qty	ID	name	price	Qty				
1	C2	...	...	4	B1	...	...	7	D2	...	...				
2	D1	...	...	5	A3	...	...	8	B2	...	...				
3	A2	...	...	6	C1	...	...	9	A1	...	...				

# Logical Order vs Physical Order

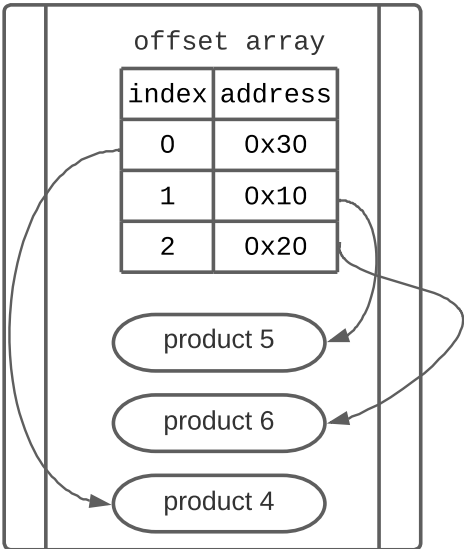
when insert unsorted data to clustered index table and all rows will be stored on the same page Like :

```
Insert into [dbo].[Product] Values ( 5 , 'Product 5' , 100 , 20 );
Insert into [dbo].[Product] Values ( 6, 'Product 6' , 120, 17);
Insert into [dbo].[Product] Values ( 4, 'Product 4' , 50, 26);
```

you think the data will be stored with physical order in the same page like :

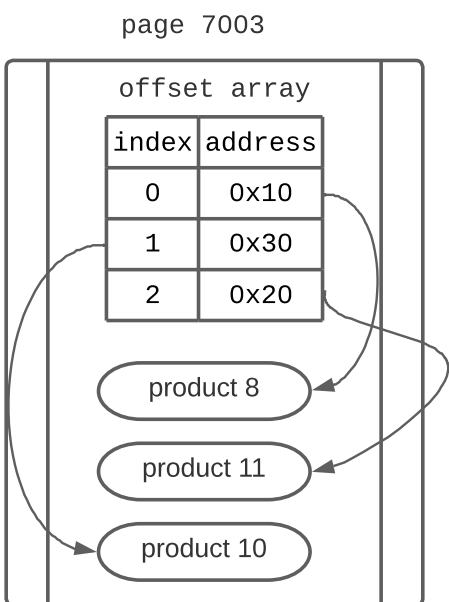


but actually database engine will stored them as they insert first , then will sort or logically order them by offset array and the data will be appear as they ordered physically.



page splits

assume we have this page

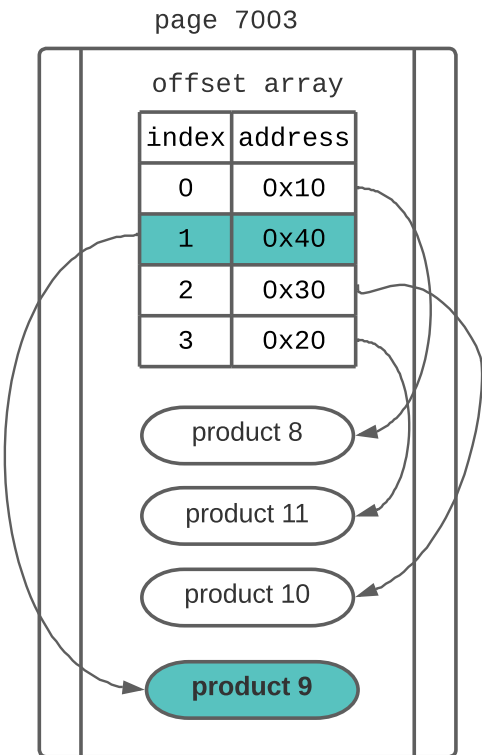


```
Insert into [dbo].[Product] Values ( 9 , 'Product 9' , 170 , 90 );
```

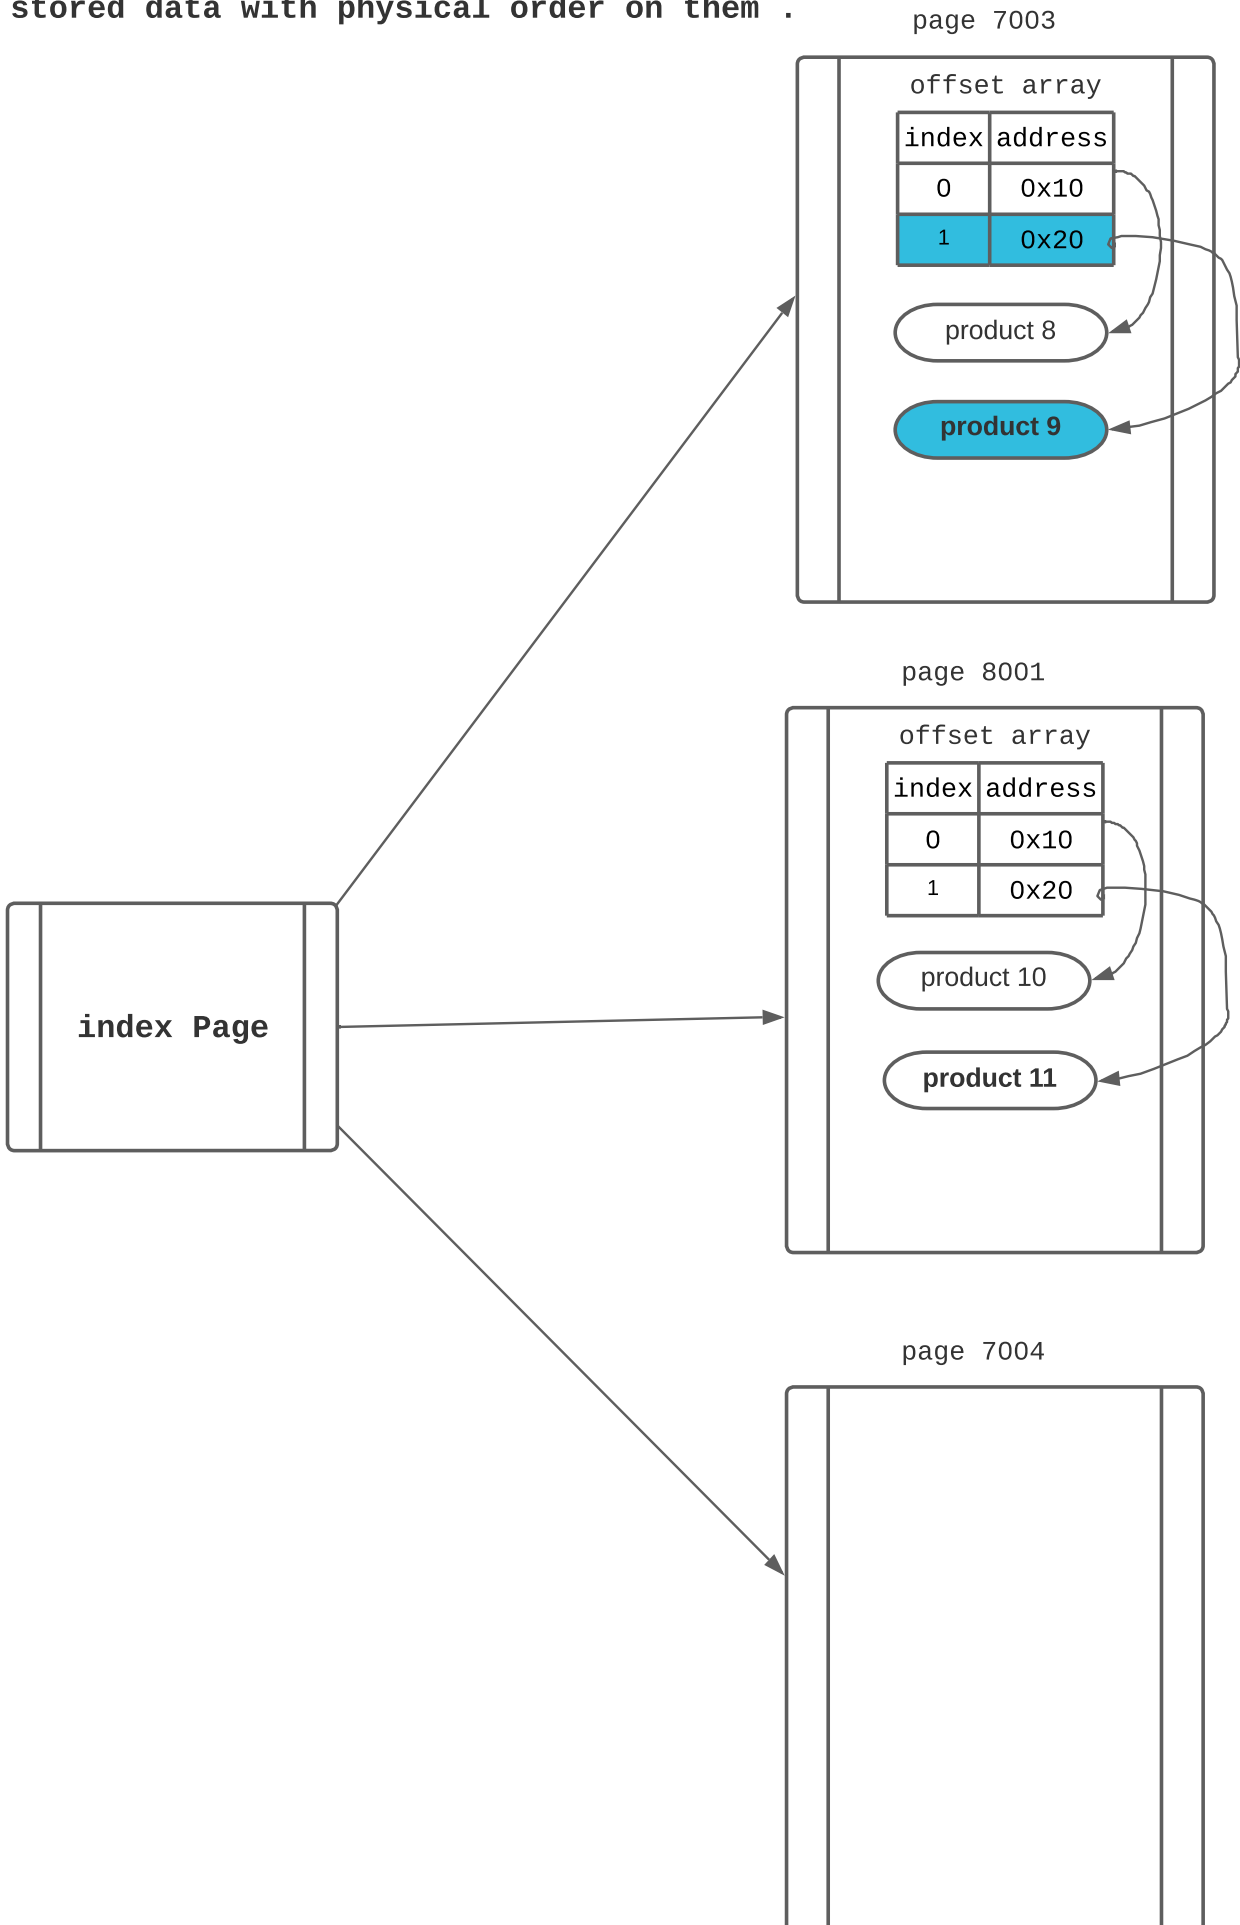
how database engine will insert this row physically ?

first step :  
database engine will **determine** the data page which the row must be inserted to it to keep the order , then database engine will **check** if there available space to insert this row at this page ?

second step :  
**if yes** : database engine will **attach** this row at the end of data page and **refactor** offset array to maintenance order of data



third step :  
**if no** : database engine will **split** the page to two pages and stored data with physical order on them .





## Clustered Index Seek

```
graph LR; A[Clustered Index Seek] --- B[algorithm ( binary search )]; A --- C[The Engine Read only Required Pages or Rows Directly From Disk Or Buffer]; A --- D[Select * from Product Where Id = 5]; D --- E["in clustered index table where id has clustered index :  
- engine read only the page that contain row with id = 5 from disk or buffer.  
- then return it to the client query"]
```

algorithm ( binary search )

The Engine Read only Required Pages or Rows Directly From Disk Or Buffer

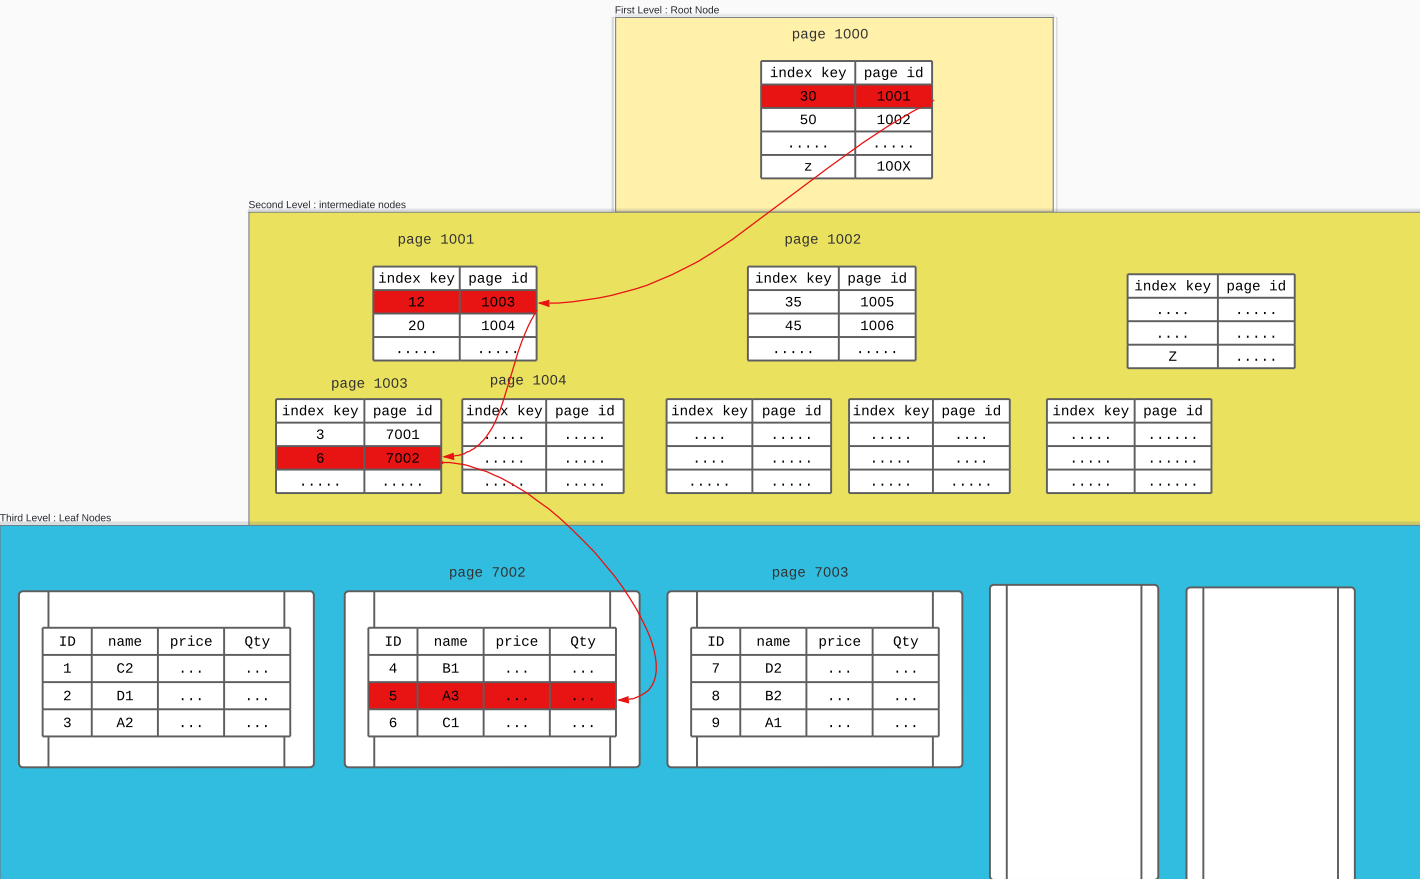
Select \* from Product Where Id = 5

in clustered index table where id has clustered index :

- engine read only the page that contain row with id = 5 from disk or buffer.
- then return it to the client query

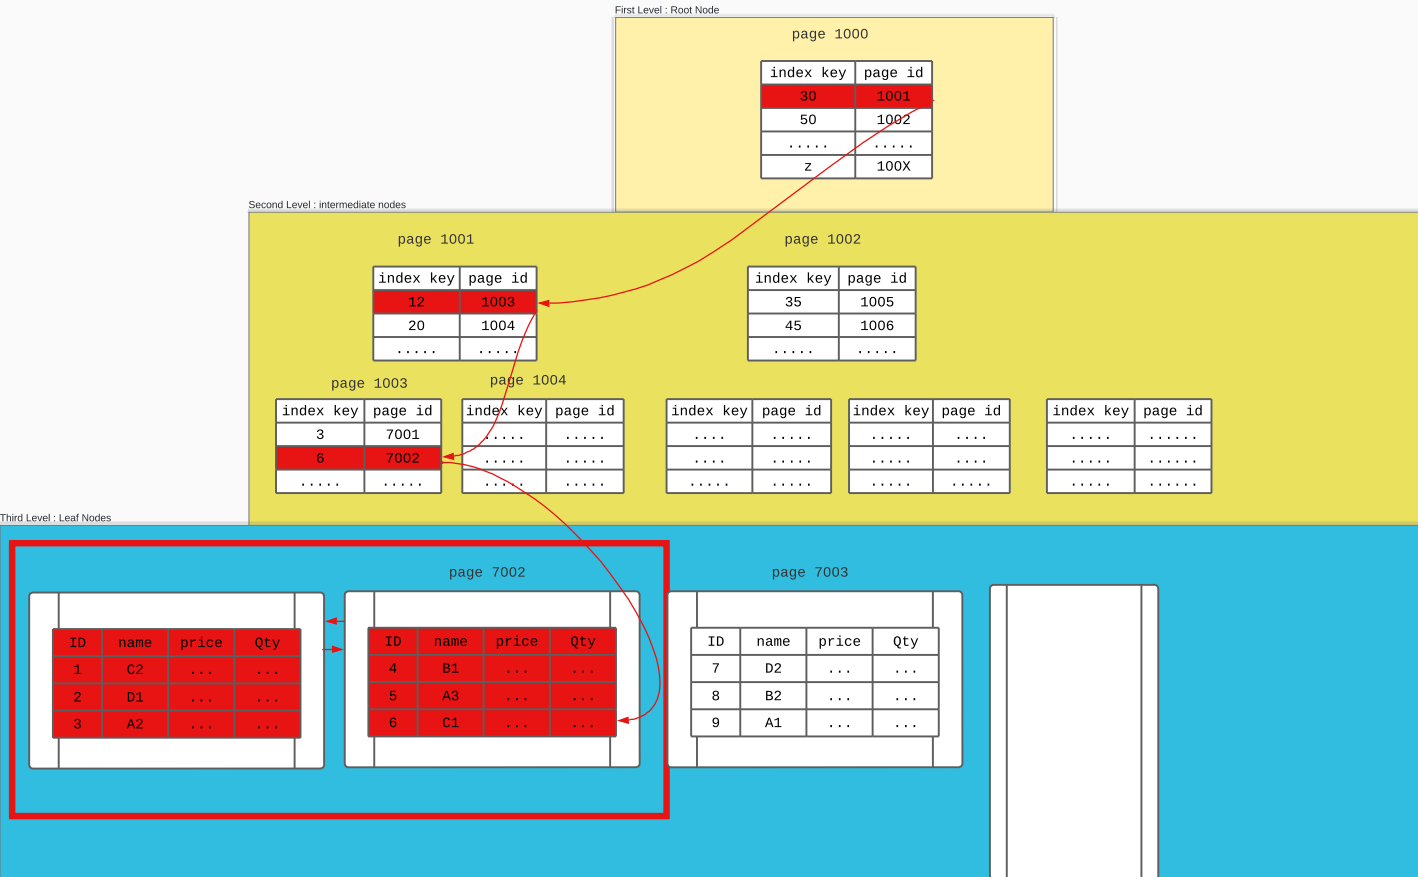
# Clustered Index Seek Example : SELECT \* FROM Product WHERE ID = 5

## Clustered Index On ID Column



# Clustered Index Seek Example : SELECT \* FROM Product WHERE ID <= 6

## Clustered Index On ID Column



# Clustered Index Seek Complexity

Algorithm : Binary Search

Complexity ~=  
number of required indexes pages (b-tree height )  
+  
number of required data pages ( actual data pages in last level )

why page is the unit ?  
page is the basic unit for the database engine not row.  
if the page contain multiple row and we want only to  
return one row , the engine will read the whole page  
from disk and filter on database server memory.

for ex :  
Select \* from product where  
id = 5

number of indexes pages : 3

number of actual data page : 1

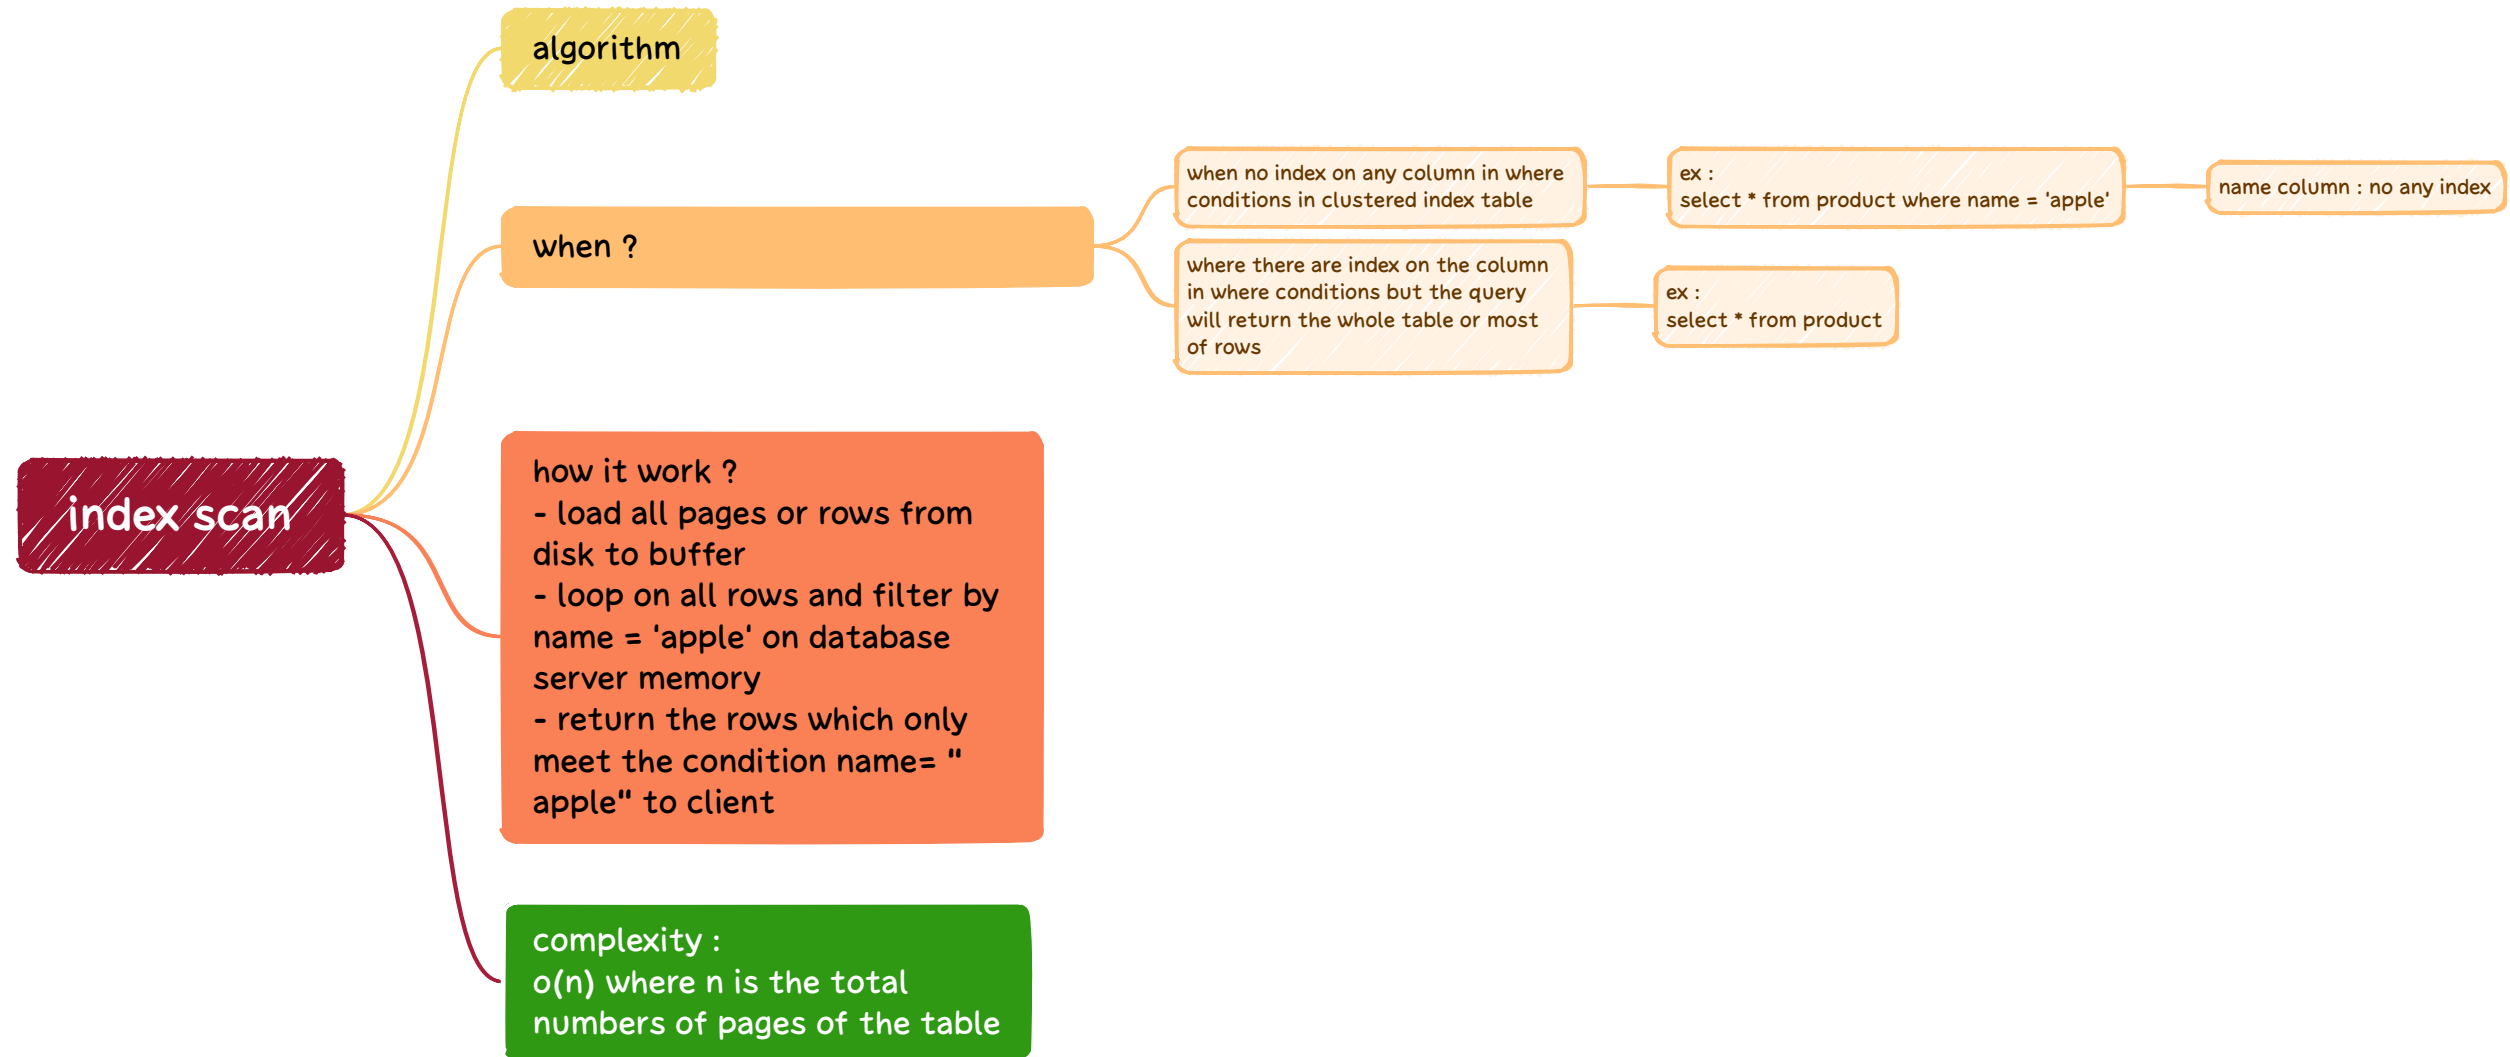
total : 4 pages

for ex :  
select \* from product where  
id < 7

number of indexes pages : 3

number of actual data page : 2

total : 5 pages



SQLQuery1.sql - DESKTOP-GNFE269\MSSQLSERVER01\Indexes\_Tutorials (DESKTOP-GNFE269\lenovo (51)) - Microsoft SQL Server Management Studio

File Edit View Query Project Tools Window Help

Indexes\_Tutorials

Execute

Object Explorer

Connect

DESKTOP-GNFE269\MSSQLSERVER01 (C)

Databases

System Databases

Database Snapshots

Indexes\_Tutorials

Database Diagrams

Tables

System Tables

FileTables

External Tables

Graph Tables

dbo.Product

Views

External Resources

Synonyms

Programmability

Service Broker

Storage

Security

Server Objects

Replication

PolyBase

Always On High Availability

Management

Integration Services Catalogs

SQL Server Agent (Agent XPs disabled)

XEvent Profiler

SQLQuery1.sql - DE...FE269\lenovo (51))

```
SELECT [ID]
      ,[Name]
      ,[Price]
      ,[AvailableQuantity]
FROM [Indexes_Tutorials].[dbo].[Product]
WHERE NAME = 'nisl Cras'
```

100 %

Results Messages Execution plan

	ID	Name	Price	AvailableQuantity
14	880526	nisl Cras	4073	99
15	977228	nisl Cras	661	11
16	994695	nisl Cras	946	47
17	114637	nisl Cras	8062	73
18	174995	nisl Cras	2096	73
19	180667	nisl Cras	5489	87
20	184273	nisl Cras	7025	89
21	354898	nisl Cras	707	81
22	476680	nisl Cras	8481	61
23	479832	nisl Cras	7527	23
24	530622	nisl Cras	8727	65
25	531883	nisl Cras	3458	34
26	573425	nisl Cras	4621	99
27	584657	nisl Cras	1663	51
28	780774	nisl Cras	5848	54
29	898454	nisl Cras	7667	27
30	911902	nisl Cras	1009	82
31	913730	nisl Cras	3633	45
32	949284	nisl Cras	3314	25

Activate Windows

Query executed successfully.

DESKTOP-GNFE269\MSSQLSERVER... DESKTOP-GNFE269\lenovo... Indexes\_Tutorials 00:00:00 32 rows

Ln 1 Col 1

SQLQuery1.sql - DESKTOP-GNFE269\MSSQLSERVER01\Indexes\_Tutorials (DESKTOP-GNFE269\lenovo (51)) - Microsoft SQL Server Management Studio

File Edit View Query Project Tools Window Help

Object Explorer

Connect +

DESKTOP-GNFE269\MSSQLSERVER01 (C)

Databases

System Databases

Database Snapshots

Indexes\_Tutorials

Database Diagrams

Tables

System Tables

FileTables

External Tables

Graph Tables

dbo.Product

Views

External Resources

Synonyms

Programmability

Service Broker

Storage

Security

Server Objects

Replication

PolyBase

Always On High Availability

Management

Integration Services Catalogs

SQL Server Agent (Agent XPs disabled)

XEvent Profiler

SQLQuery1.sql - DE...FE269\lenovo (51)\*

SELECT [ID]  
[Name]  
[Price]  
[AvailableQuantity]  
FROM [Indexes\_Tutorials].[dbo].[Product]  
WHERE NAME = 'nisl Cras'

Query 1: Query cost (relative to the batch):  
SELECT [ID], [Name], [Price], [AvailableQuantity]  
Missing Index (Impact 99.4348): CREATE NONCLUSTERED INDEX [IX\_Product\_ID] ON [Indexes\_Tutorials].[dbo].[Product] ([ID])

Execution plan

Parallelism (Gather Streams) ← Clustered Index Scan (Clustered)

Cost: 5 %  
0.081s  
\$2 of  
29 (110%)

Clustered Index Scan (Clustered)

Scanning a clustered index, entirely or only a range.

Physical Operation	Clustered Index Scan
Logical Operation	Clustered Index Scan
Actual Execution Mode	Row
Estimated Execution Mode	Row
Storage	RowStore
Number of Rows Read	1000006
Actual Number of Rows for All Executions	32
Actual Number of Batches	0
Estimated I/O Cost	4.92312
Estimated Operator Cost	5.47321 (95%)
Estimated Subtree Cost	5.47321
Estimated CPU Cost	0.550082
Estimated Number of Executions	1
Number of Executions	2
Estimated Number of Rows for All Executions	28.9683
Estimated Number of Rows Per Execution	28.9683
Estimated Number of Rows to be Read	1000010
Estimated Row Size	49 B
Actual Rebinds	0
Actual Rewinds	0
Ordered	False
Node ID	1

Predicate  
[Indexes\_Tutorials].[dbo].[Product].[Name]='nisl Cras'

Object  
[Indexes\_Tutorials].[dbo].[Product].[IX\_Product\_ID]

Output List  
[Indexes\_Tutorials].[dbo].[Product].[ID] [Indexes\_Tutorials].[dbo].[Product].[Name] [Indexes\_Tutorials].[dbo].[Product].[Price] [Indexes\_Tutorials].[dbo].[Product].[AvailableQuantity]

Query executed successfully.

Activate Windows

27°C 27° سماء صافية

6:19 PM 9/22/2022

SQLQuery1.sql - DESKTOP-GNFE269\MSSQLSERVER01\Indexes\_Tutorials (DESKTOP-GNFE269\lenovo (51)) - Microsoft SQL Server Management Studio

File Edit View Query Project Tools Window Help

Object Explorer

- DESKTOP-GNFE269\MSSQLSERVER01 (C)
- Databases
  - System Databases
  - Database Snapshots
  - Indexes\_Tutorials
    - Database Diagrams
    - Tables
      - System Tables
      - FileTables
      - External Tables
      - Graph Tables
      - dbo.Product
    - Views
    - External Resources
    - Synonyms
    - Programmability
    - Service Broker
    - Storage
    - Security
  - Security
  - Server Objects
  - Replication
  - PolyBase
  - Always On High Availability
  - Management
  - Integration Services Catalogs
  - SQL Server Agent (Agent XPs disabled)
  - XEvent Profiler

SQLQuery1.sql - DE...FE269\lenovo (51))

```
SELECT [ID]
      , [Name]
      , [Price]
      , [AvailableQuantity]
FROM [Indexes_Tutorials].[dbo].[Product]
WHERE NAME = 'nisi1 Cras'
```

100 %

Results Messages Execution plan

Query 1: Query cost (relative to the batch): 100%

SELECT [ID], [Name], [Price], [AvailableQuantity] FROM [Indexes\_Tutorials].[dbo].[Product] WHERE ([NAME]=8)

Missing Index (Impact 99.4348): CREATE NONCLUSTERED INDEX [<Name of Missing Index, sysname,>] ON [dbo].[Product] ([Name])

Execution Plan:

- SELECT (Cost: 0 %)
- Parallelism (Gather Streams) (Cost: 5 %)
- Clustered Index Scan (Clustered) [Product].[IX\_Product\_ID] (Cost: 95 %)

Query executed successfully.

DESKTOP-GNFE269\MSSQLSERVER... DESKTOP-GNFE269\lenovo... Indexes\_Tutorials 00:00:00 // 32 rows

Activate Windows

Ready

27°C سماء صافية 6:34 PM 9/22/2022 ENG



# non-Clustered Index

in-disk b-tree data structure consist of 3 levels :

1- root node : contain  
- non-clustered index keys  
- with address of pages of other lower index pages

2- intermediate nodes : contain  
- non-clustered index keys  
- with address of pages of other lower index pages

3- leaf nodes : contain  
- pointers to actual data

Nonclustered indexes have a structure separate from the data rows or actual table

Indexes Key Is Sorted in Indexes Pages in Root or intermediate nodes

can used with heap table or clustered table

pointers :

in case clustered table : clustered index keys

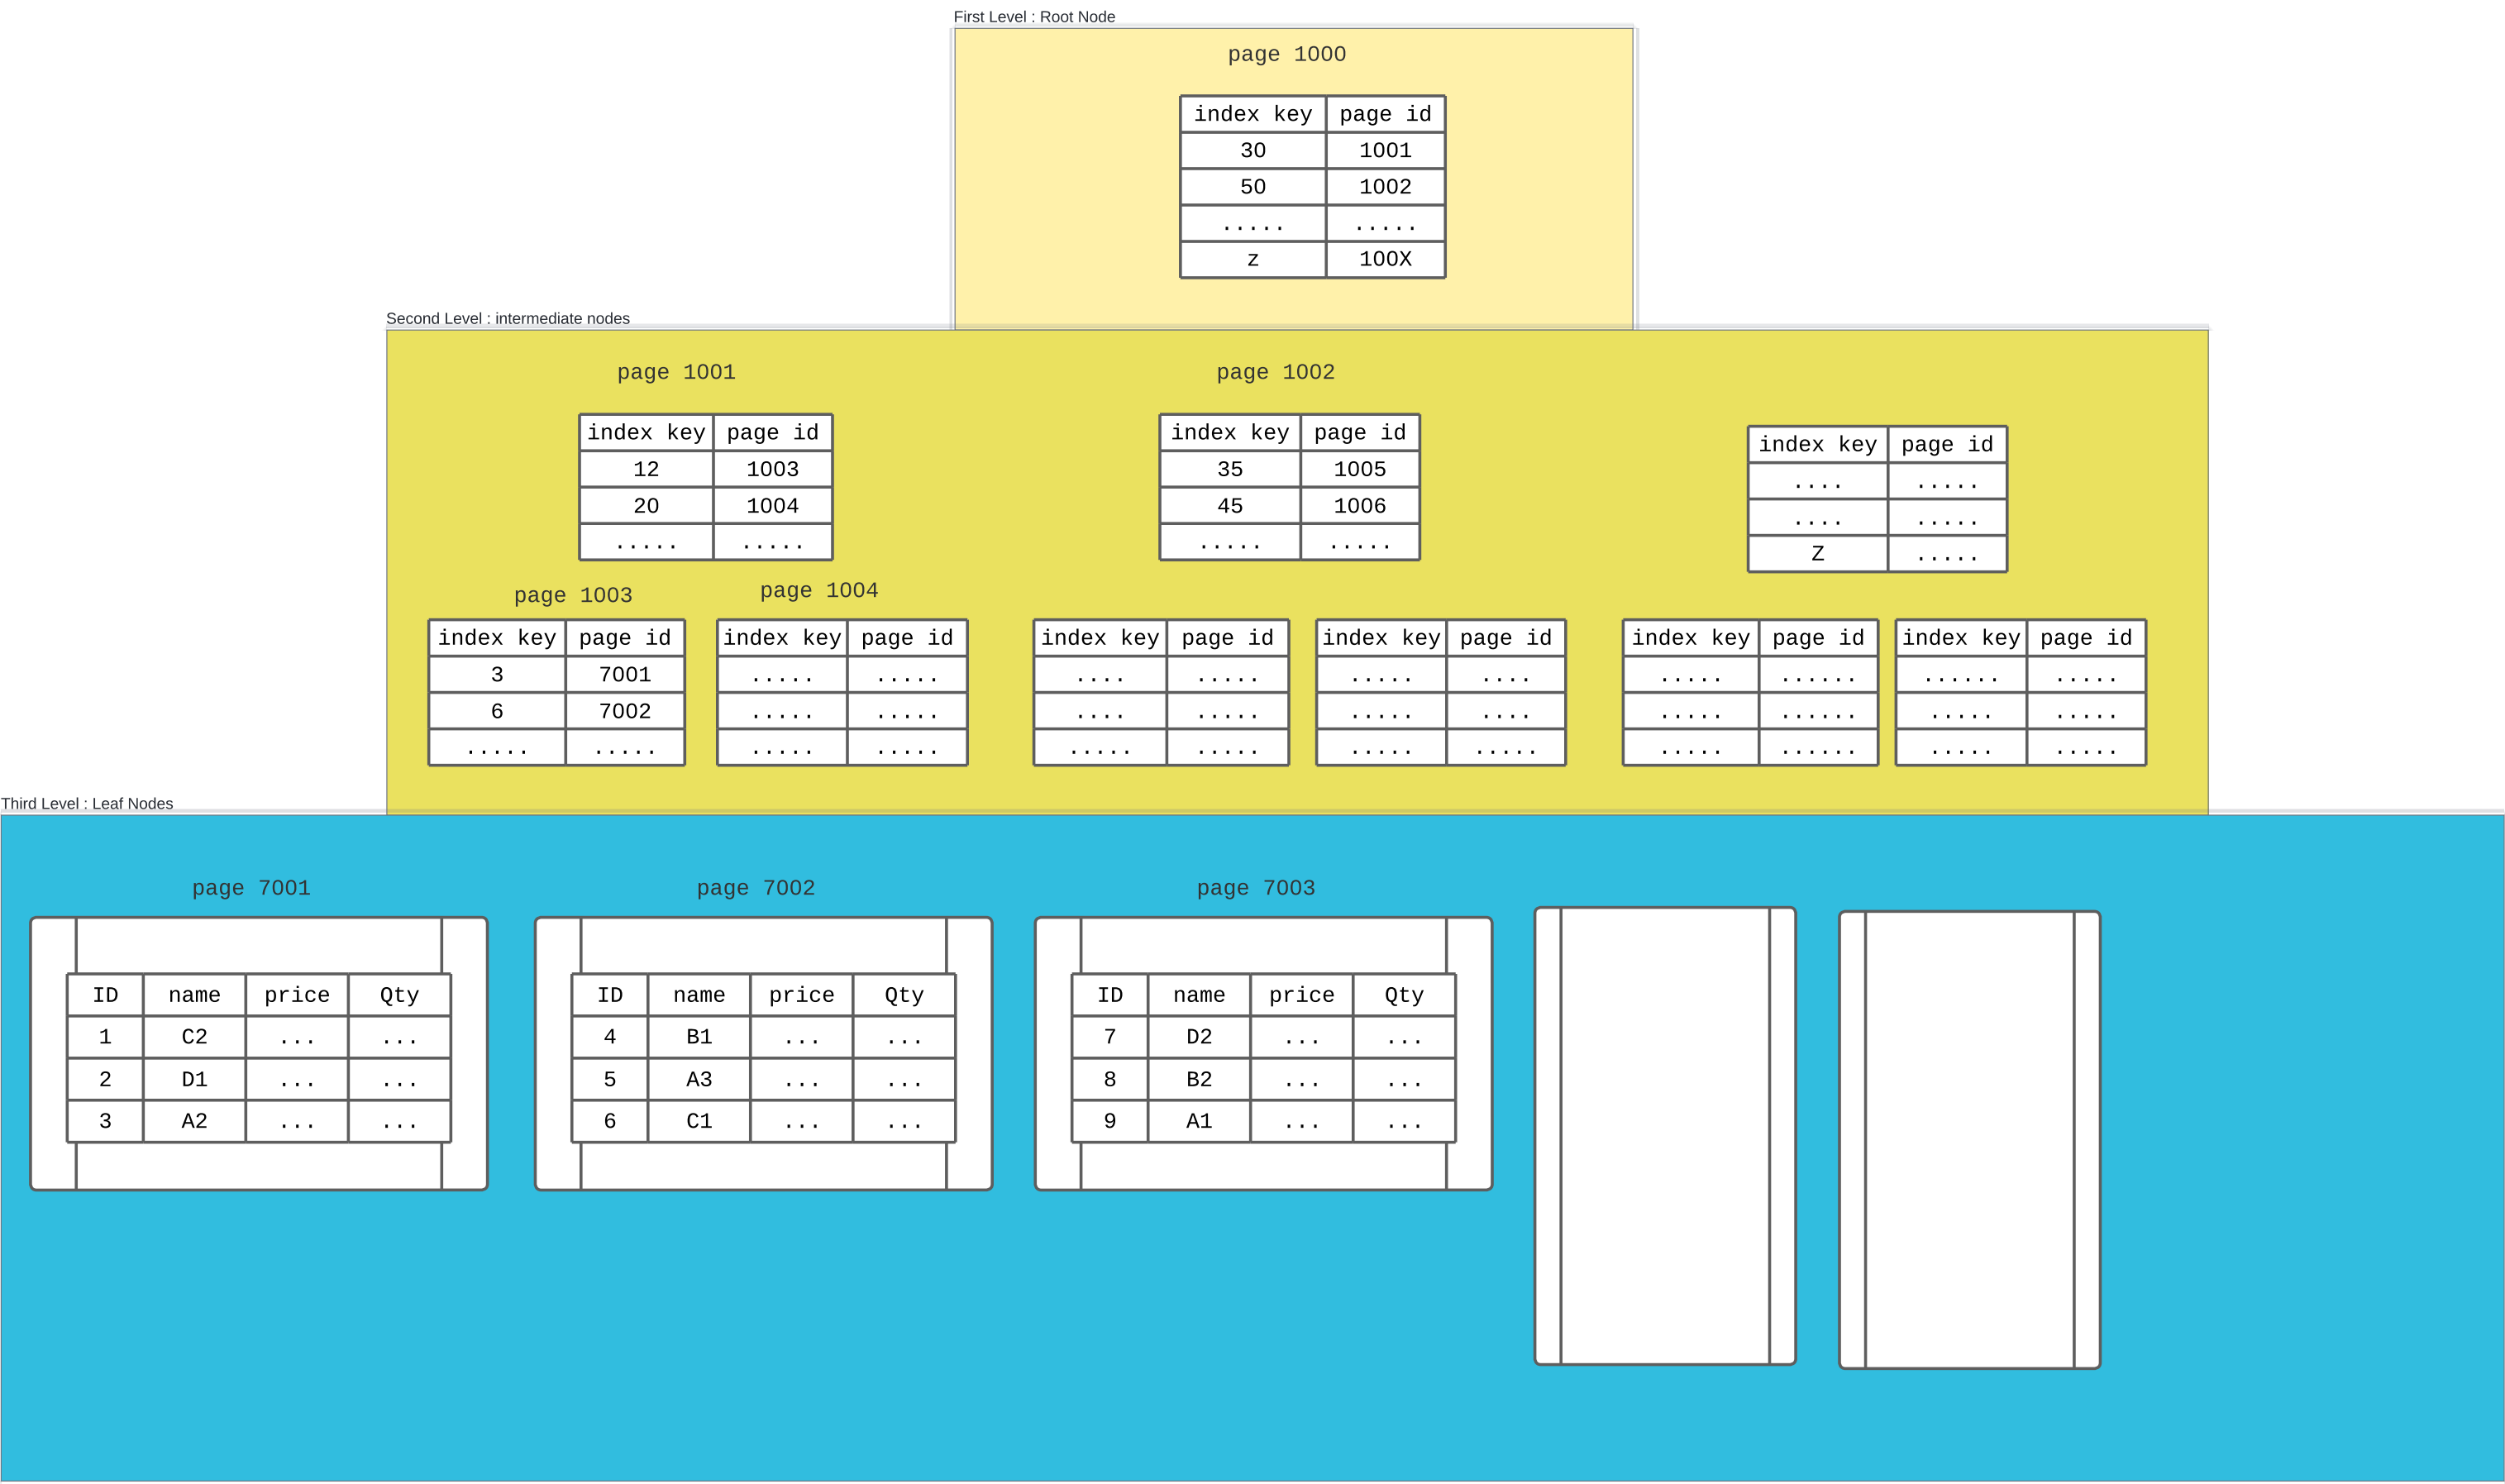
in case heap table :  
ROWID (File + Page + row )

table can contain many non-clustered indexes

Create Non-Clustered Index on  
Name Column on Product Table

```
CREATE NonCLUSTERED  
INDEX [IX_Product_Name]  
ON [dbo].[Product]  
(  
    [Name] Asc  
)
```

Clustered Index On ID Column



Non Clustered Index On Name Column



SQLQuery1.sql - DESKTOP-GNFE269\MSSQLSERVER01\Indexes\_Tutorials (DESKTOP-GNFE269\lenovo (62)) - Microsoft SQL Server Management Studio

File Edit View Query Project Tools Window Help

Object Explorer

- DESKTOP-GNFE269\MSSQLSERVER01 (C)
- Databases
  - System Databases
  - Database Snapshots
  - Indexes\_Tutorials
    - Database Diagrams
    - Tables
      - System Tables
      - FileTables
      - External Tables
      - Graph Tables
      - dbo.Product
    - Views
    - External Resources
    - Synonyms
    - Programmability
    - Service Broker
    - Storage
    - Security
  - Server Objects
  - Replication
  - PolyBase
  - Always On High Availability
  - Management
  - Integration Services Catalogs
  - SQL Server Agent (Agent XPs disabled)
  - XEvent Profiler

SQLQuery1.sql - DE...FE269\lenovo (62)\*

```
CREATE NONCLUSTERED INDEX [IX_Product_Name]
ON [dbo].[Product]
(
    [Name] Asc
)
SELECT [ID]
, [Name]
, [Price]
, [AvailableQuantity]
FROM [Indexes_Tutorials].[dbo].[Product]
WHERE NAME = 'nisl Cra'
```

100 %

Results Messages Live Query Statistics Execution plan

Estimated query cost (relative to the batch): 100%

progress:100% SELECT [ID] , [Name] , [Price] , [AvailableQuantity] FROM [Indexes\_Tutorials].[dbo].[Product] WHERE NAME = 'nisl Cra'

Execution plan diagram:

- SELECT (32 of 29 (11.0%))
- Nested Loops (Inner Join) (32 of 29 (11.0%))
- Index Seek (NonClustered) ([Product].[IX\_Product\_Name]) (32 of 29 (11.0%))
- Key Lookup (Clustered) ([Product].[IX\_Product\_ID]) (32 of 29 (11.0%))

Query executed successfully.

DESKTOP-GNFE269\MSSQLSERVER... DESKTOP-GNFE269\lenovo... Indexes\_Tutorials 00:00:00 36 rows

Ready Ln 12 Col 25 Ch 26 INS

24°C 9:46 PM 9/23/2022

SQLQuery1.sql - DESKTOP-GNFE269\MSSQLSERVER01\Indexes\_Tutorials (DESKTOP-GNFE269\lenovo (62)) - Microsoft SQL Server Management Studio

Object Explorer

- DESKTOP-GNFE269\MSSQLSERVER01 (C)
- Databases
  - System Databases
  - Database Snapshots
  - Indexes\_Tutorials
    - Database Diagrams
    - Tables
      - System Tables
      - FileTables
      - External Tables
      - Graph Tables
      - dbo.Product
    - Views
    - External Resources
    - Synonyms
    - Programmability
    - Service Broker
    - Storage
    - Security
  - Server Objects
  - Replication
  - PolyBase
  - Always On High Availability
  - Management
  - Integration Services Catalogs
  - SQL Server Agent (Agent XPs disabled)
  - XEvent Profiler

SQLQuery1.sql - DE...FE269\lenovo (62)\*

```
CREATE NONCLUSTERED INDEX [IX_Product_Name]
ON [dbo].[Product]
(
    [Name] Asc
)

SELECT [ID]
, [Name]
, [Price]
, [AvailableQuantity]
FROM [Indexes_Tutorials].[dbo].[Product]
WHERE NAME = 'nisl Cras'
```

Estimated query progress: 100%

Query 1: Query cost (relative) progress: 100%

SELECT [ID], [Name], [Price]

100 %

SELECT

Nested Loops (Inner Join) 32 of 29 (11.0%)

Index Seek (Product)

Key Lookups (Product)

Query executed successfully.

Index Seek (NonClustered)

Scan a particular range of rows from a nonclustered index.

Estimated operator progress: 100%

Physical Operation	Index Seek
Logical Operation	Index Seek
Estimated Execution Mode	Row
Storage	RowStore
Actual Number of Rows for All Executions	32
Estimated Operator Cost	0.0033135 (4%)
Estimated I/O Cost	0.003125
Estimated Subtree Cost	0.0033135
Estimated CPU Cost	0.0001885
Number of Executions	1
Estimated Number of Executions	1
Estimated Number of Rows Per Execution	28.604
Estimated Number of Rows for All Executions	28.604
Estimated Number of Rows to be Read	28.604
Estimated Row Size	45 B
Ordered	True
Node ID	2

Object

[Indexes\_Tutorials].[dbo].[Product] [IX\_Product\_Name]

Output List

Uniq1001, [Indexes\_Tutorials].[dbo].[Product].ID, [Indexes\_Tutorials].[dbo].[Product].Name

Seek Predicates

Seek Keys[1]: Prefix: [Indexes\_Tutorials].[dbo].[Product].Name = 'nisl Cras'

Scalar Operator('nisl Cras')

TOP-GNFE269\lenovo... Indexes\_Tutorials 00:00:00 36 rows

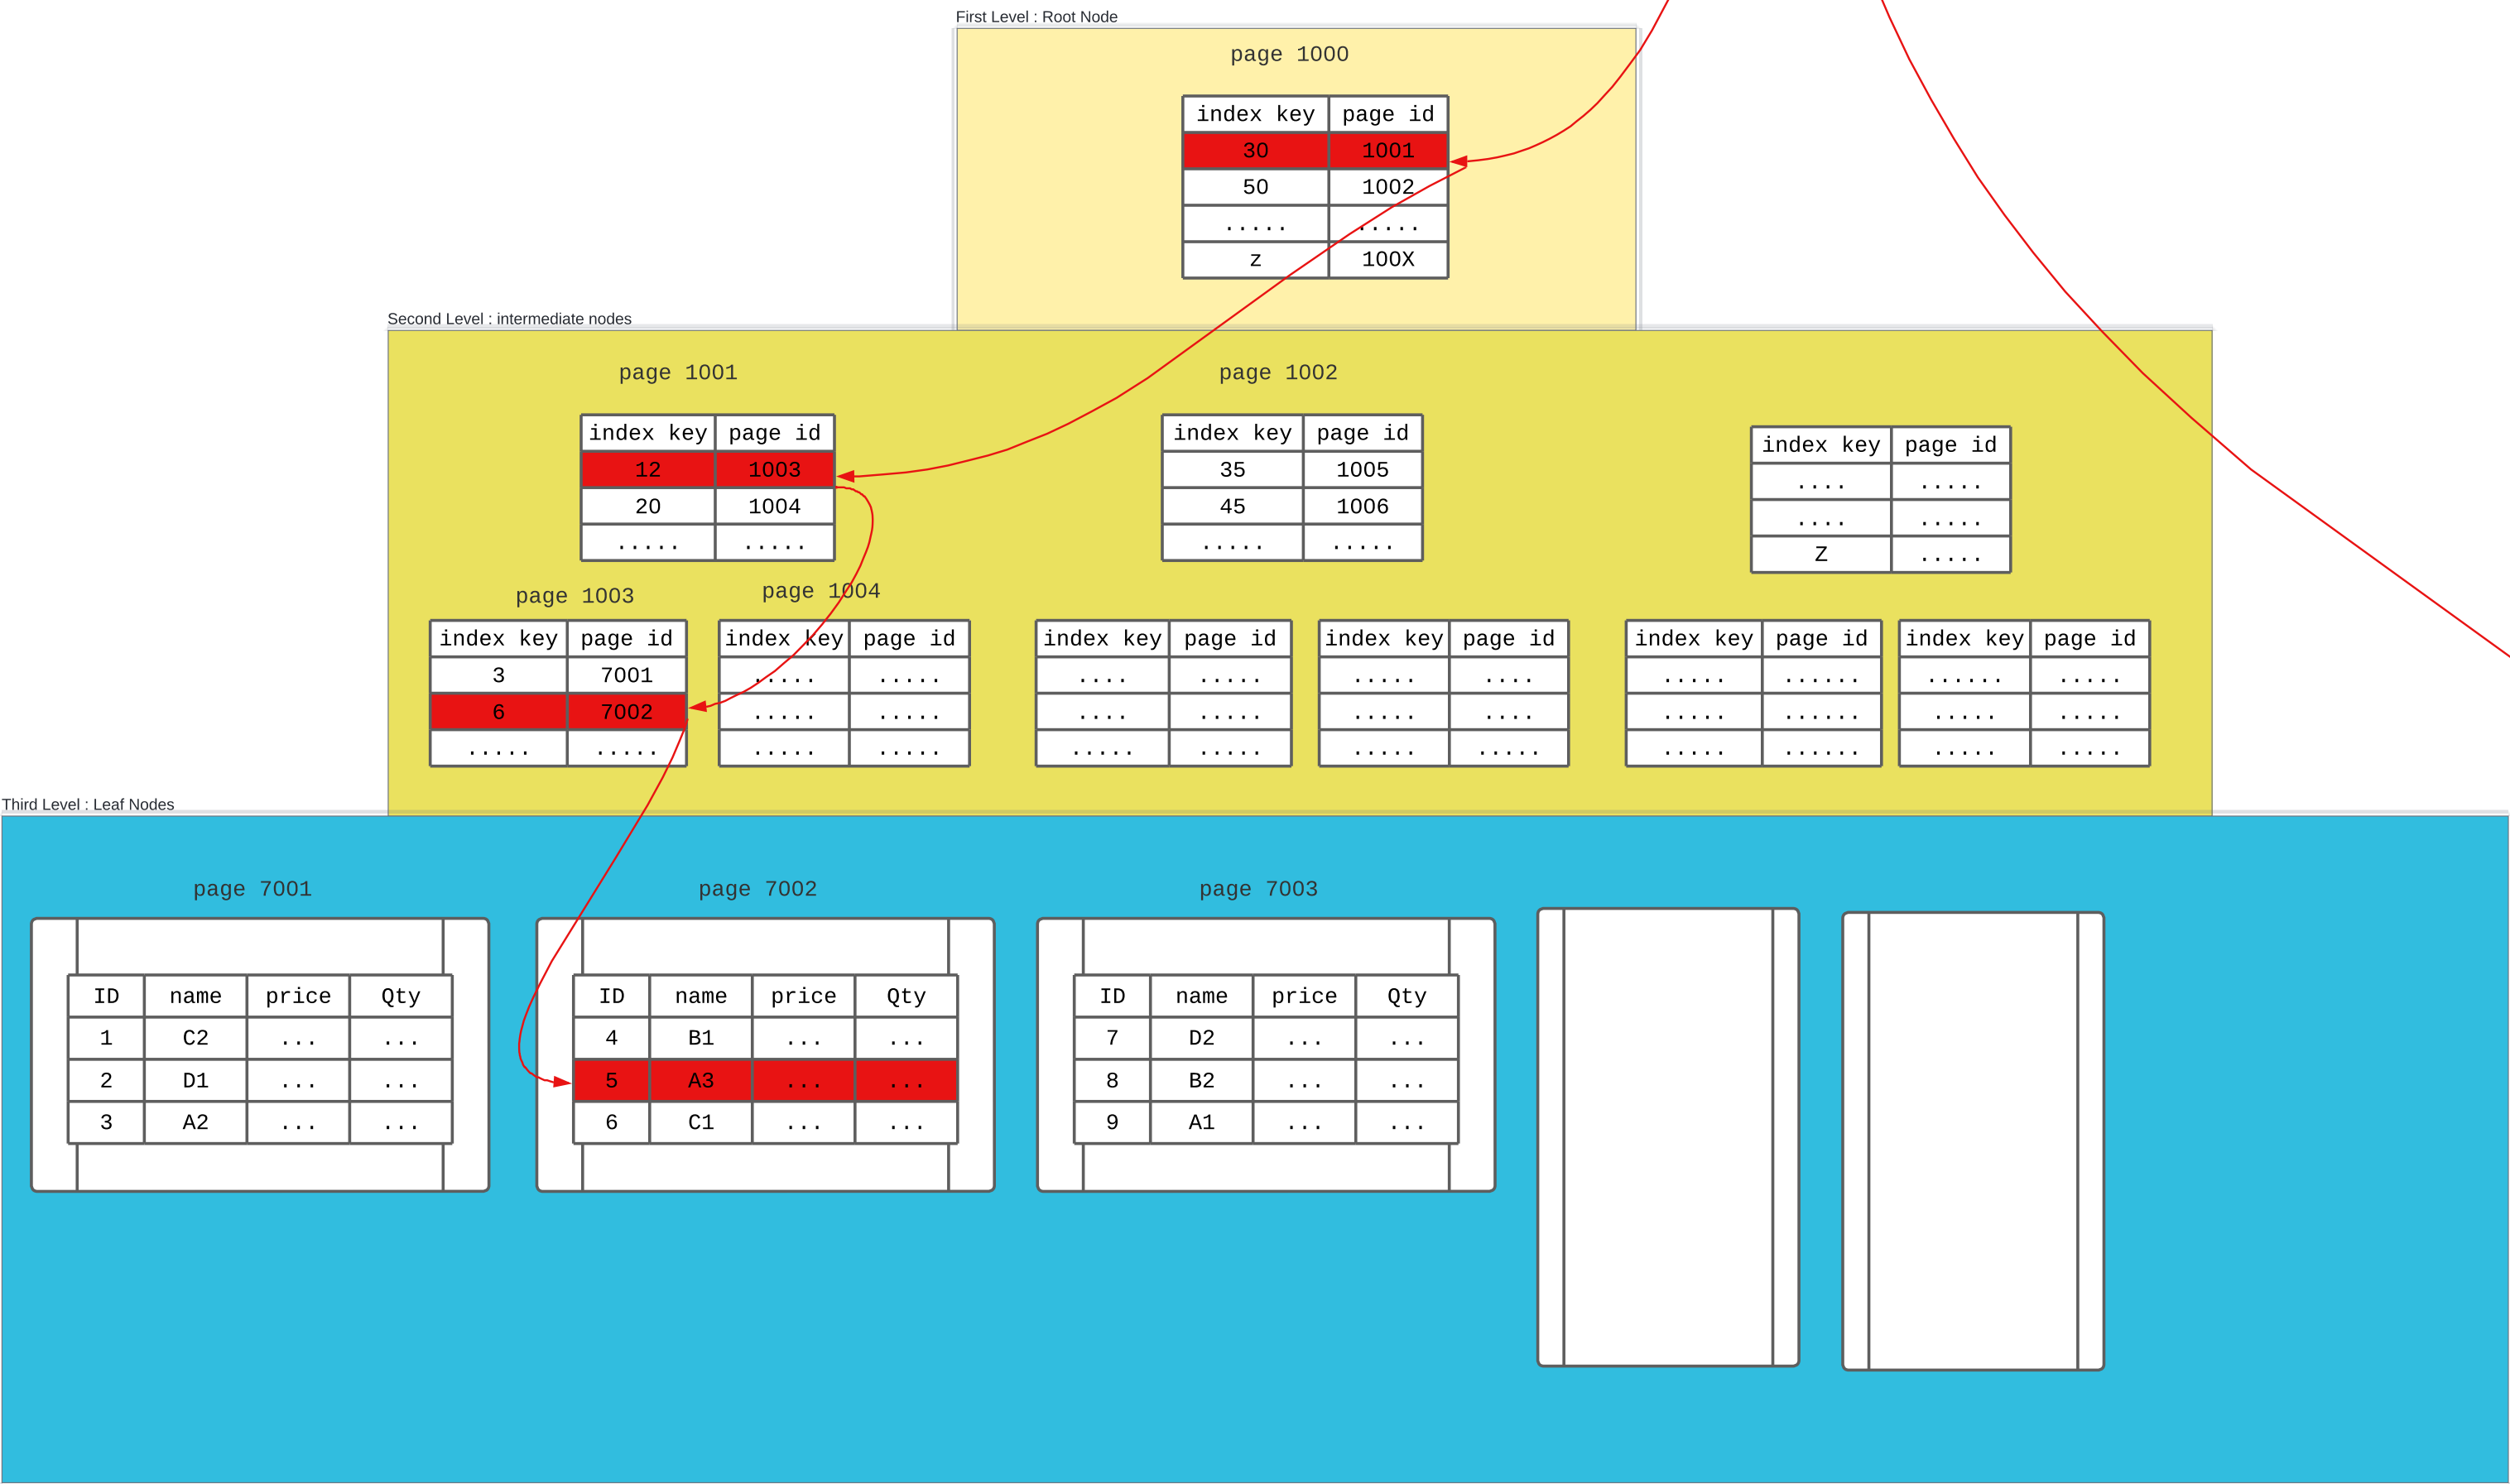
Ready

24°C 9:47 PM 9/23/2022

SELECT \* FROM Product where Name = 'A3'

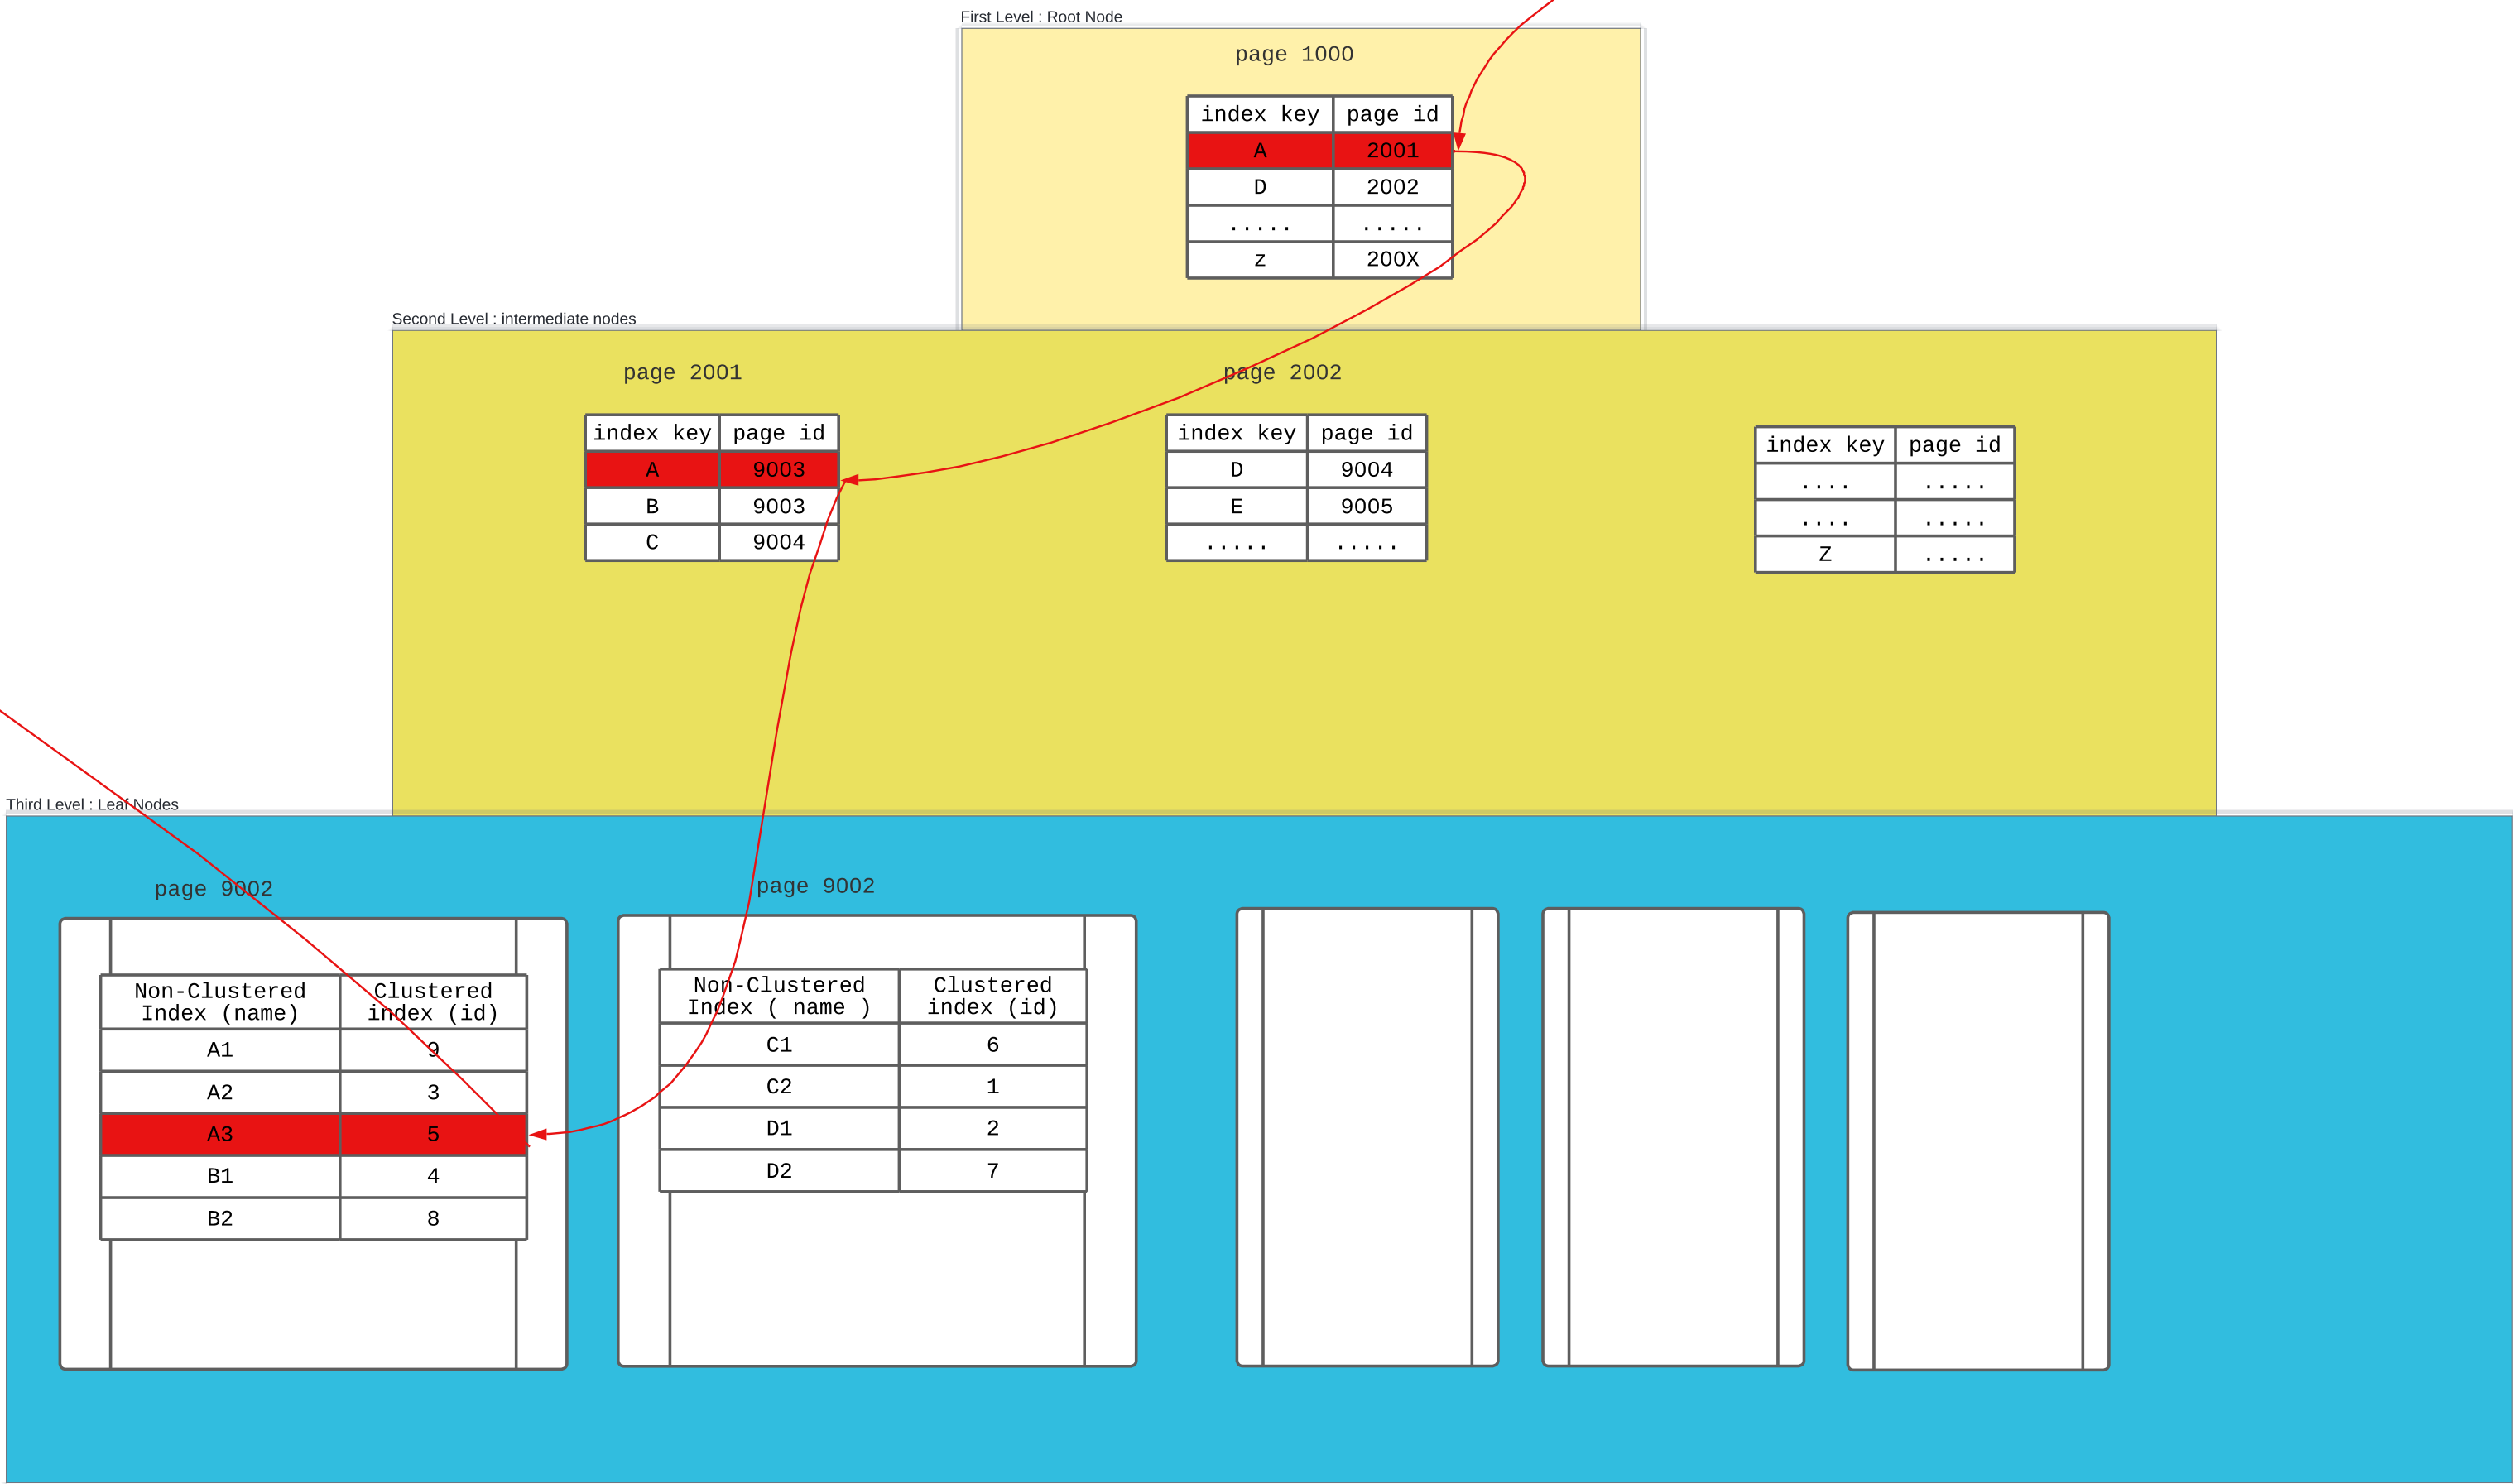
2- Key Lookup on Clustered Index  
Search about Product Has ID = 5

Clustered Index On ID Column



1- Non Clustered Index Seek On Name Column  
Get id of Product Has Name = A3

Non Clustered Index On Name Column



## non-clustered index with include

```
graph LR; A[non-clustered index with include] --- B[if our table has a lot of fields, but most business cases require to read only a few fields or columns]; A --- C[use include to prevent key lookup operation by add copy of the selected columns in non-clustered index]; A --- D[this called covered query: where all the columns in the select query return from the non-clustered indexes];
```

if our table has a lot of fields, but most business cases require to read only a few fields or columns

use include to prevent key lookup operation by add copy of the selected columns in non-clustered index

this called covered query: where all the columns in the select query return from the non-clustered indexes

## create non-clustered index with include example

ex :  
select name , price from  
product

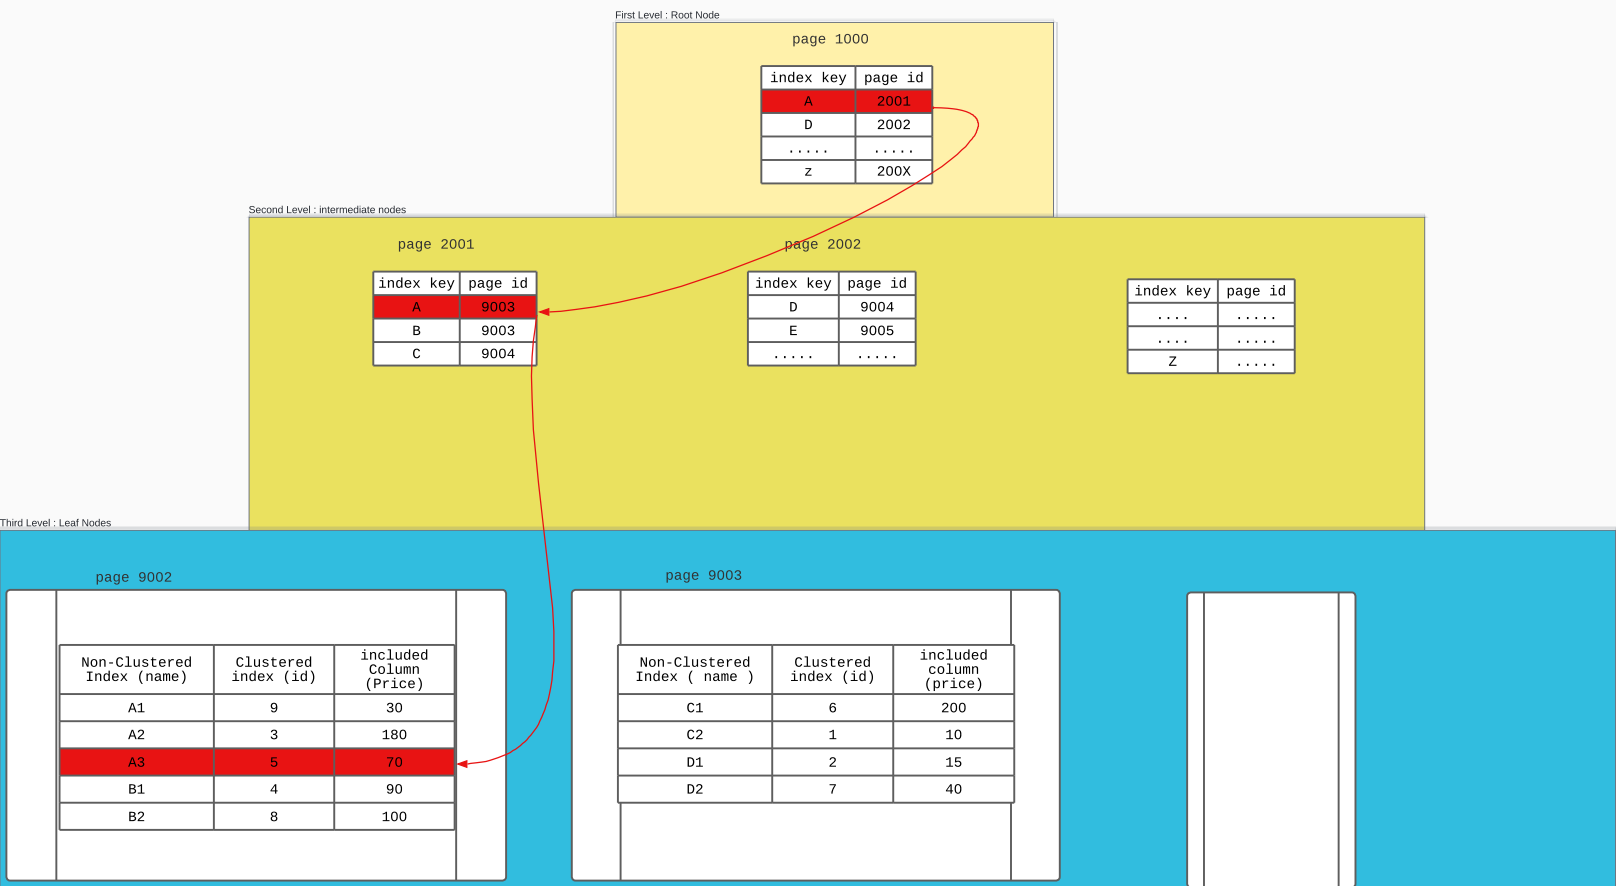
```
CREATE NonCLUSTERED  
INDEX [IX_Product_Name]  
ON [dbo].[Product]  
(  
    [Name] Asc  
)  
Include (Price)
```



SELECT NAME , PRICE FROM Product  
Where Name = 'A3'

1- Non Clustered Index Seek

Non Clustered Index On Name Column  
with include ( Price Column )



SQLQuery1.sql - DESKTOP-GNFE269\MSSQLSERVER01\Indexes\_Tutorials (DESKTOP-GNFE269\lenovo (59)) - Microsoft SQL Server Management Studio

File Edit View Query Project Tools Window Help

Object Explorer

Connect +

TOP-GNFE269\MSSQLSERVER01 (SQL Serve ^

databases

- System Databases
- Database Snapshots
- Indexes\_Tutorials
  - Database Diagrams
  - Tables
    - System Tables
    - FileTables
    - External Tables
    - Graph Tables
    - dbo.Product
      - Columns
      - Keys
      - Constraints
      - Triggers
      - Indexes
        - IX\_Product\_ID (Clustered)
      - Statistics
  - Views
  - External Resources
  - Synonyms
  - Programmability
  - Service Broker
  - Storage
  - Security
- Server Objects
  - Replication

SQLQuery1.sql - DE...FE269\lenovo (59))

```
CREATE NONCLUSTERED INDEX [IX_Product_Name]
ON [dbo].[Product]
(
    [Name] Asc
)
Include (Price)

SELECT [Name]
,[Price]
FROM [Indexes_Tutorials].[dbo].[Product]
WHERE NAME = 'nisl Cras'
```

100 %

Results Messages Execution plan

Query 1: Query cost (relative to the batch): 100%

SELECT [Name],[Price] FROM [Indexes\_Tutorials].[dbo].[Product] WHERE [NAME]=@1

Index Seek (NonClustered)

Index: [Product].[IX\_Product\_Name]

Cost: 100 %

0.000s

32 of

29 (110%)

SELECT

Cost: 0 %

Query executed successfully.

DESKTOP-GNFE269\MSSQLSERVER... DESKTOP-GNFE269\lenovo... Indexes\_Tutorials 00:00:01 32 rows

Ready

31°C 2:03 PM 9/24/2022

